

Files in the Dataset

1. yelp_academic_dataset_business.json

- **1.2M+ businesses**
- **Key fields:**
 - business_id: Unique ID
 - name: Business name
 - address, city, state, postal_code, latitude, longitude
 - stars: Average star rating (float)
 - review_count: Number of reviews
 - is_open: Open status (1 = open)
 - categories: Comma-separated categories (e.g., "Japanese, Sushi Bars, Restaurants")
 - attributes: Nested dict with metadata (e.g., "WiFi", "Parking", "Good for kids")

Use in RL:

- To enrich item context (e.g., categories, popularity)
 - To support action selection features
-

2. yelp_academic_dataset_review.json

- **8.6M+ reviews**
- **Key fields:**
 - review_id
 - user_id, business_id: Relational links
 - stars: Actual rating (1–5 stars)
 - date: Timestamp of the review
 - text: Freeform review text (not always used in structured models)
 - useful, funny, cool: Count of feedback from other users

Use in RL:

- Forms the core of (state, action, reward) triplets

- The rating becomes reward
 - Enables sorting by time to form sequences
-

3. yelp_academic_dataset_user.json

- **2.1M+ users**
- **Key fields:**
 - user_id, name
 - review_count
 - yelping_since: Registration date
 - friends: Comma-separated user_ids
 - useful, funny, cool: Aggregate votes
 - fans, elite: Popularity/social influence indicators
 - average_stars: Avg. rating given

Use in RL:

- Used to enrich the **state** with user profile (average_stars, review_count, friend count)
 - Can personalize action-value function
-

4. yelp_academic_dataset_checkin.json

- **Check-in records per business**
- **Key fields:**
 - business_id
 - date: Comma-separated timestamps (e.g., "2018-01-01 09:00:00, 2018-01-02 18:00:00")

Use in RL:

- Extract time-based popularity (morning / afternoon / evening visits)
 - Serves as context feature for action (business popularity signal)
-

5. yelp_academic_dataset_tip.json

- **Short tips (mini reviews)**
- **Key fields:**
 - user_id, business_id, text, date
 - compliment_count

Use in RL:

- Optional text-based feedback signal
- Can complement reviews for sentiment analysis or explanation

Schema Relationships

User (user_id)



Review (user_id, business_id, stars, date)



Business (business_id, categories, attributes)



Check-in (business_id, date)

In RL Framework

RL Concept	Dataset Source	Explanation
state	Review history	Sequence of previous business_id, stars, date, user profile, business features
action	Business	The next business to recommend
reward	Review	Usually 1 if stars ≥ 4 , else 0
next_state	Review	User's updated sequence after action
context_features	Business, Check-in, User	Additional info like categories, check-in frequency, average_stars

Component	Feature Name	Source File	Description
state	business_id	review.json	Previously interacted businesses
	stars	review.json	Rating the user gave to that business
	date	review.json	Timestamp of the review (for sequence)
	time_segment	derived from date	Morning / Afternoon / Evening
	days_since_action	computed from date difference	Used for time-decay weight
	recency_weight	computed	$\exp(-\lambda \times \text{time_diff})$
	user_profile.review_count	user.json	Number of reviews the user wrote
	user_profile.average_stars	user.json	User's mean rating
	user_profile.friend_count	user.json	Proxy for social influence
	business_checkin.morning/afternoon/evening	checkin.json	Popularity at different times
	business_categories	business.json	Category tags (multi-label)
action	business_id	from state or model selection	The business recommended to the user
	action_time_segment	review.json	When the recommendation occurred
	business_checkin	checkin.json	Contextual popularity
	business_categories	business.json	Contextual metadata
reward	reward	review.json	Binary reward: 1 if stars \geq 4, else 0
next_state	same as state shifted forward	review.json	Resulting state after the user reacts to the action

Notes:

- state is a **sequence** of 5 (or more) past interactions
- action is typically a **single business_id**

- reward is derived from the user's rating
- $\text{next_state} = \text{state} + \text{action}$